

## Practical-1

AIM: Program on Prolog

precious(gold).

precious(silver).

female(mary).

father(surya,ramesh).

likes(surya,food).

likes(surya,car).

likes(surya,bike).

likes(surya,book).

likes(shreya,chocolate).

likes(sudha,food).

likes(surya,coffee).

likes(samudra,chocolate).

likes(samanta,X) :- likes(X,chocolate).

```
SWI-Prolog (AMD64, Multi-threaded, version 9.0.4)
File Edit Settings Run Debug Help
Welcome to SWI-Prolog (threaded, 64 bits, version 9.0.4)
SWI-Prolog comes with ABSOLUTELY NO WARRANTY. This is free software.
Please run ?- license. for legal details.

For online help and background, visit https://www.swi-prolog.org
For built-in help, use ?- help(Topic). or ?- apropos(Word).

?-
% c:/users/lab2_51/documents/pract1 compiled 0.00 sec, -2 clauses
?- precious(gold)
|
ERROR: Stream user_input:9:15 Syntax error: Unexpected end of file
?- precious(gold).
true.

?- precious(silver).
true.

?- precious(platinum).
false.

?- likes(surya,food).
true
Unknown action: ( (h for help)
Action?
ERROR: Type error: 'character_code' expected, found '-1' (an integer)
ERROR: In:
ERROR: [11] char_code(_14626,-1)
ERROR: [10] '$in_reply'(-1,'?h') at c:/program files/swipl/boot/init.pl:1037
?- likes(shreya,book).
false.

?- likes(samanta,X)
|
ERROR: Stream user_input:33:0 Syntax error: Unexpected end of file
?- likes(samanta,X).
X = shreya .

?- likes(samanta,X).
X = shreya ■
```

## Practical-2

AIM: Implementation of water jug problem using prolog

```
water_jug(X,Y):- X>4,Y<3,write('4L jug overflow. '),nl.
water_jug(X,Y):- X<4,Y>3,write('3L jug overflow. '),nl.
water_jug(X,Y):- X>4,Y>3,write('Both jugs overflow. '),nl.

water_jug(0, 0) :- write('4L:0 & 3L:0'), nl, water_jug(4, 0).

water_jug(4, 0) :- write('4L:4 & 3L:0 (Action: Fill 4L jug.)'), nl,
water_jug(1, 3).

water_jug(1, 3) :- write('4L:1 & 3L:3 (Action: Pour water from 4L to
3L jug.)'), nl, water_jug(1, 0).

water_jug(1, 0) :- write('4L:1 & 3L:0 (Action: Empty 3L jug.)'), nl,
water_jug(0, 1).

water_jug(0, 1) :- write('4L:0 & 3L:1 (Action: Pour water from 4L jug
to 3L jug.)'), nl, water_jug(4, 1).

water_jug(4, 1) :- write('4L:4 & 3L:1 (Action: Fill 4L jug.)'), nl,
water_jug(2, 3).

water_jug(2, 3) :- write('4L:2 & 3L:3 (Action: Pour water from 4L to
3L jug untill 3L jug is full.)'), nl, water_jug(2, 0).

water_jug(2, 0) :- write('4L:2 & 3L:0 (Action: Empty 3L jug. Goal
State reached..)'), nl.
```

```
Welcome to SWI-Prolog (threaded, 64 bits, version 9.3.3-21-g1fbd70c21)
SWI-Prolog comes with ABSOLUTELY NO WARRANTY. This is free software.
Please run ?- license. for legal details.
```

```
For online help and background, visit https://www.swi-prolog.org
For built-in help, use ?- help(Topic). or ?- apropos(Word).
```

```
?-
% c:/Users/USER/Downloads/demoBhushan.pl compiled 0.00 sec, 11 clauses
```

```
?- waterjug(9,0).
Correct to: "water_jug(9,0)"? yes
4L jug overflow.
true.
```

```
?- waterjug(0,6).
Correct to: "water_jug(0,6)"? yes
3L jug overflow.
true.
```

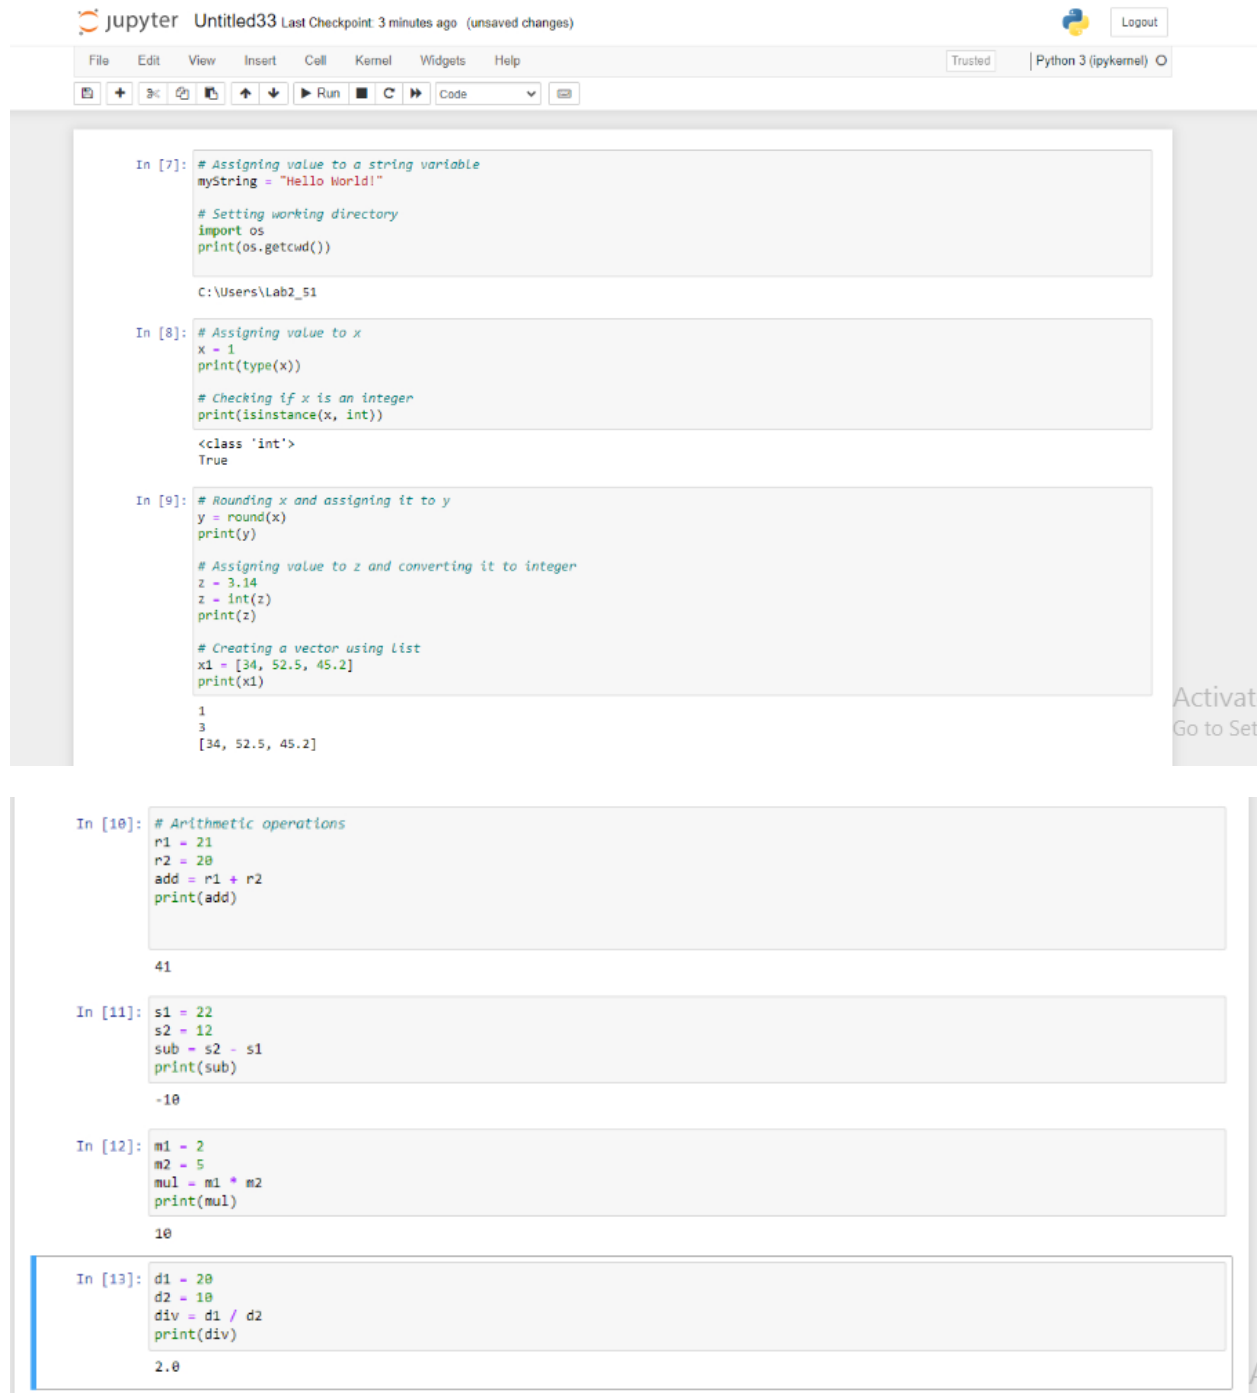
```
?- waterjug(6,6).
Correct to: "water_jug(6,6)"? yes
Both jugs overflow.
true.
```

```
?-
|   waterjug(0,0).
Correct to: "water_jug(0,0)"? yes
4L:0 & 3L:0
4L:4 & 3L:0 (Action: Fill 4L jug.)
4L:1 & 3L:3 (Action: Pour water from 4L to 3L jug.)
4L:1 & 3L:0 (Action: Empty 3L jug.)
4L:0 & 3L:1 (Action: Pour water from 4L jug to 3L jug.)
4L:4 & 3L:1 (Action: Fill 4L jug.)
4L:2 & 3L:3 (Action: Pour water from 4L to 3L jug untill 3L jug is full.)
4L:2 & 3L:0 (Action: Empty 3L jug. Goal State reached..)
true.
```

```
?- ■
```

## Practical 3

AIM: Introduction to Python Programming



The image shows a Jupyter Notebook interface with the following content:

**Jupyter Untitled33** Last Checkpoint: 3 minutes ago (unsaved changes) Python 3 (ipykernel)

**In [7]:**

```
# Assigning value to a string variable
myString = "Hello World!"

# Setting working directory
import os
print(os.getcwd())
```

C:\Users\Lab2\_51

**In [8]:**

```
# Assigning value to x
x = 1
print(type(x))

# Checking if x is an integer
print(isinstance(x, int))
```

<class 'int'>  
True

**In [9]:**

```
# Rounding x and assigning it to y
y = round(x)
print(y)

# Assigning value to z and converting it to integer
z = 3.14
z = int(z)
print(z)

# Creating a vector using list
x1 = [34, 52.5, 45.2]
print(x1)
```

1  
3  
[34, 52.5, 45.2]

**In [10]:**

```
# Arithmetic operations
r1 = 21
r2 = 20
add = r1 + r2
print(add)
```

41

**In [11]:**

```
s1 = 22
s2 = 12
sub = s2 - s1
print(sub)
```

-10

**In [12]:**

```
m1 = 2
m2 = 5
mul = m1 * m2
print(mul)
```

10

**In [13]:**

```
d1 = 20
d2 = 10
div = d1 / d2
print(div)
```

2.0

## Practical-4

AIM: Introduction to Python Libraries

```
In [1]: import numpy as np
import pandas as pd
```

```
In [2]: df = pd.read_csv('C:\\Users\\USER\\Desktop\\eda\\mtcars.csv')
```

```
In [3]: print(df)
```

	model	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	\
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	
3	Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	
5	Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	
6	Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	
7	Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	
8	Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	
9	Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	
10	Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	
11	Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	
12	Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	
13	Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	
14	Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	
15	Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	
16	Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	
17	Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	
18	Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	
19	Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	
20	Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	
21	Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	
22	AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	
23	Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	
24	Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	
25	Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	
26	Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	
27	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	
28	Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	
29	Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	
30	Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	
31	Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	

	gear	carb
0	4	4
1	4	4
2	4	1
3	3	1
4	3	2
5	3	1
6	3	4
7	4	2
8	4	2
9	4	4
10	4	4
11	3	3
12	3	3

```
In [5]: df.head()
```

Out[5]:

	model	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
0	Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
1	Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
2	Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
3	Hornet 4 Drive	21.4	6	259.0	110	3.08	3.215	19.44	1	0	3	1
4	Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
5	Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1

```
In [ ]: df.head(10)
```

```
In [6]: df.info()
```

```

class 'pandas.core.frame.DataFrame'>
RangeIndex: 32 entries, 0 to 31
Data columns (total 12 columns):
#   Column      Non-Null Count  Dtype
---  -
0   model       32 non-null    object
1   mpg         32 non-null    float64
2   cyl         32 non-null    int64
3   disp       32 non-null    float64
4   hp         32 non-null    int64
5   drat       32 non-null    float64
6   wt         32 non-null    float64
7   qsec       32 non-null    float64
8   vs         32 non-null    int64
9   am         32 non-null    int64
10  gear       32 non-null    int64
11  carb       32 non-null    int64
dtypes: float64(5), int64(6), object(1)
memory usage: 3.1+ KB

```

```
In [9]: df.isnull()
```

Out[9]:

[illegible]

In [10]: `df.isnull().sum()`

Out[10]:

```
model      0
mpg        0
cyl        0
disp       0
hp         0
drat       0
wt         0
qsec       0
vs         0
am         0
gear       0
carb       0
dtype: int64
```

In [11]: `df.tail()`

Out[11]:

	model	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
27	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.9	1	1	5	2
28	Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.5	0	1	5	4
29	Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.5	0	1	5	6
30	Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.6	0	1	5	8
31	Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.6	1	1	4	2

In [12]: `df.describe()`

Out[12]:

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
count	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000	32.000000
mean	20.090625	6.187500	230.721875	146.687500	3.596563	3.217250	17.848750	0.437500	0.406250	3.687500	2.812500
std	6.026948	1.785922	123.938694	68.562868	0.534679	0.978457	1.786943	0.504016	0.498991	0.737804	1.615000
min	10.400000	4.000000	71.100000	52.000000	2.760000	1.513000	14.500000	0.000000	0.000000	3.000000	1.000000
25%	15.425000	4.000000	120.825000	96.500000	3.080000	2.581250	16.892500	0.000000	0.000000	3.000000	2.000000
50%	19.200000	6.000000	196.300000	123.000000	3.695000	3.325000	17.710000	0.000000	0.000000	4.000000	2.000000
75%	22.800000	8.000000	326.000000	180.000000	3.920000	3.610000	18.900000	1.000000	1.000000	4.000000	4.000000
max	33.900000	8.000000	472.000000	335.000000	4.930000	5.424000	22.900000	1.000000	1.000000	5.000000	8.000000

In [13]: `df.size`

Out[13]: 384

In [14]: `df.shape`

Out[14]: (32, 12)

In [15]: `df.ndim`

Out[15]: 2

In [16]: `df.at[4,'model']`

Out[16]: 'Hornet Sportabout'

In [17]: `df.at[4,'disp']`

Out[17]: 360.0

In [18]: `df.iat[3,4]`

Out[18]: 110

In [20]: `df.loc[:, 'model']`



```
Out[20]: 0      Mazda RX4
1      Mazda RX4 Wag
2      Datsun 710
3      Hornet 4 Drive
4      Hornet Sportabout
5      Valiant
6      Duster 360
7      Merc 240D
8      Merc 230
9      Merc 280
10     Merc 280C
11     Merc 450SE
12     Merc 450SL
13     Merc 450SLC
14     Cadillac Fleetwood
15     Lincoln Continental
16     Chrysler Imperial
17     Fiat 128
18     Honda Civic
19     Toyota Corolla
20     Toyota Corona
21     Dodge Challenger
22     AMC Javelin
23     Camaro Z28
24     Pontiac Firebird
25     Fiat X1-9
26     Porsche 914-2
27     Lotus Europa
28     Ford Pantera L
29     Ferrari Dino
30     Maserati Bora
31     Volvo 142E
Name: model, dtype: object
```

```
In [23]: df.iloc[0:5,0:2]
```

```
Out[23]:
```

	model	mpg
0	Mazda RX4	21.0
1	Mazda RX4 Wag	21.0
2	Datsun 710	22.8
3	Hornet 4 Drive	21.4
4	Hornet Sportabout	18.7

```
In [26]: df.iloc[22:32,:]
```

```
Out[26]:
```

	model	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
22	AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
23	Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
24	Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
25	Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
26	Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
27	Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
28	Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
29	Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
30	Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
31	Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

```
In [27]: df.dtypes
```

```
Out[27]: model      object
mpg      float64
cyl      int64
disp     float64
hp       int64
drat     float64
wt       float64
qsec     float64
vs       int64
am       int64
gear     int64
carb     int64
dtype: object
```

```
In [28]: df['model'].dtype
```

```
Out[28]: dtype('O')
```

```
In [29]: df.axes
```

```
Out[29]: [RangeIndex(start=0, stop=32, step=1),  
         Index(['model', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs', 'am',  
               'gear', 'carb'],  
               dtype='object')]
```

```
In [30]: df.columns
```

```
Out[30]: Index(['model', 'mpg', 'cyl', 'disp', 'hp', 'drat', 'wt', 'qsec', 'vs', 'am',  
               'gear', 'carb'],  
               dtype='object')
```

```
In [31]: df['hp'].std()
```

```
Out[31]: 68.56286848932059
```

```
In [32]: df['mpg'].mean()
```

```
Out[32]: 20.090625000000003
```

```
In [33]: df['mpg'].median()
```

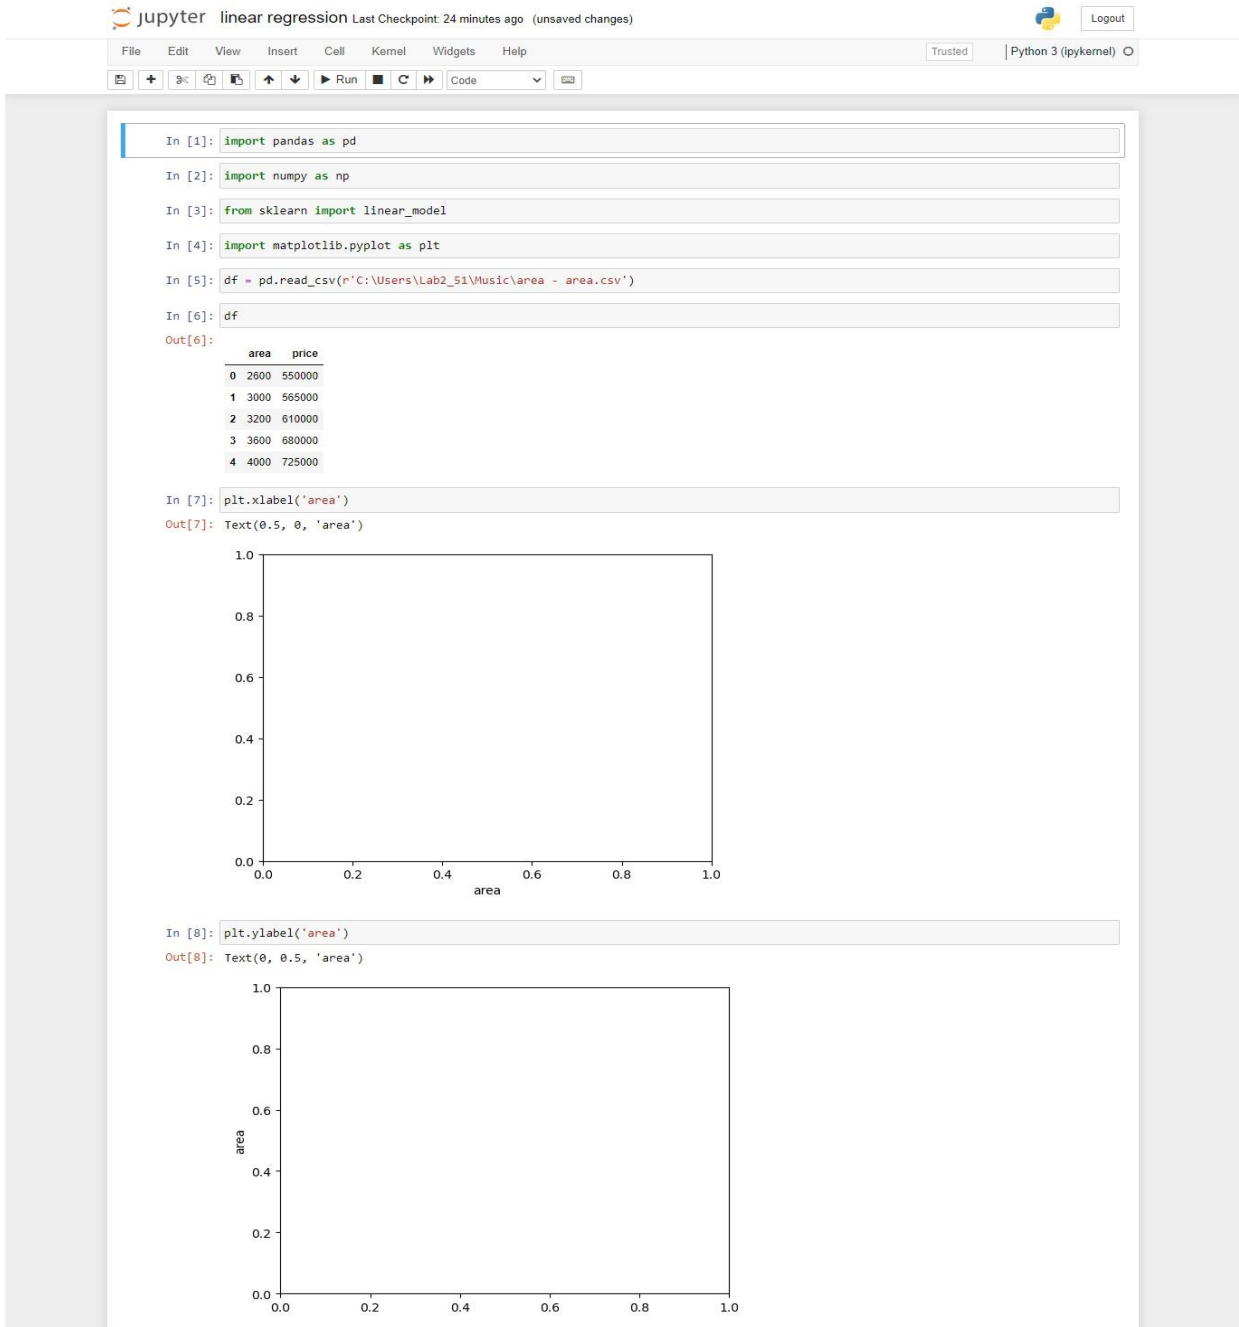
```
Out[33]: 19.2
```

```
In [34]: df['hp'].describe()
```

```
Out[34]: count      32.000000  
         mean      146.687500  
         std       68.562868  
         min       52.000000  
         25%       96.500000  
         50%      123.000000  
         75%      180.000000  
         max      335.000000  
         Name: hp, dtype: float64
```

## Practical-5

AIM: Program to implement Linear Regression



```
In [12]: price = df.price
In [13]: price
Out[13]: 0    550000
         1    565000
         2    610000
         3    680000
         4    725000
         Name: price, dtype: int64

In [14]: reg = linear_model.LinearRegression()
In [15]: reg.fit(new_df, price)
Out[15]: LinearRegression()
LinearRegression()


In [16]: reg.predict([[3300]])
C:\Users\Lab2_51\anaconda3\lib\site-packages\sklearn\base.py:420: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
  warnings.warn(
Out[16]: array([628715.75342466])

In [17]: reg.coef_
Out[17]: array([135.78767123])


In [18]: reg.intercept_
Out[18]: 180616.43835616432

In [19]: reg.predict([[5000]])
C:\Users\Lab2_51\anaconda3\lib\site-packages\sklearn\base.py:420: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
  warnings.warn(
Out[19]: array([850554.79452055])
```


```
In [21]: plt.ylabel('price',fontsize = 20)  
Out[21]: Text(0, 0.5, 'price')
```



```
In [22]: plt.scatter(df.area,df.price,color='red', marker='x')  
Out[22]: <matplotlib.collections.PathCollection at 0x1e30a01f430>
```




```
In [23]: plt.plot(df.area,reg.predict(df[['area']]),color='blue')  
Out[23]: <matplotlib.lines.Line2D at 0x1e30a09aa70>
```



```
In [25]: plt.xlabel('area',fontsize = 20)
plt.ylabel('price',fontsize = 20)
plt.scatter(df.area,df.price,color='red', marker = '+')
plt.plot(df.area,reg.predict(df[['area']]),color = 'blue')

Out[25]: [<matplotlib.lines.Line2D at 0x1e309672b30>]
```



In [ ]:

## Practical-6

AIM: Program to Implement Logistic Regression

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: credit_df = pd.read_csv('D:/Rukhsar_AIMLPacts/CreditRisk.csv')
```

```
In [3]: credit_df.shape
```

```
Out[3]: (614, 13)
```

```
In [4]: credit_df.head()
```

```
Out[4]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Pr
0	LP001002	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	

```
In [5]: credit_df.tail()
```

```
Out[5]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Pr
609	LP002978	Female	No	0	Graduate	No	2900	0.0	71	360.0	1.0	
610	LP002979	Male	Yes	3+	Graduate	No	4106	0.0	40	180.0	1.0	
611	LP002983	Male	Yes	1	Graduate	No	8072	240.0	253	360.0	1.0	
612	LP002984	Male	Yes	2	Graduate	No	7583	0.0	187	360.0	1.0	
613	LP002990	Female	No	0	Graduate	Yes	4583	0.0	133	360.0	0.0	

In [6]: credit\_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID               614 non-null    object
1   Gender                601 non-null    object
2   Married               611 non-null    object
3   Dependents            599 non-null    object
4   Education             614 non-null    object
5   Self_Employed         582 non-null    object
6   ApplicantIncome       614 non-null    int64
7   CoapplicantIncome     614 non-null    float64
8   LoanAmount            614 non-null    int64
9   Loan_Amount_Term      600 non-null    float64
10  Credit_History        564 non-null    float64
11  Property_Area         614 non-null    object
12  Loan_Status           614 non-null    int64
dtypes: float64(3), int64(3), object(7)
memory usage: 62.5+ KB
```

In [7]: credit\_df.describe()

Out[7]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status
count	614.000000	614.000000	614.000000	600.00000	564.000000	614.000000
mean	5403.459283	1621.245798	141.166124	342.00000	0.842199	0.687296
std	6109.041673	2926.248369	88.340630	65.12041	0.364878	0.463973
min	150.000000	0.000000	0.000000	12.00000	0.000000	0.000000
25%	2877.500000	0.000000	98.000000	360.00000	1.000000	0.000000
50%	3812.500000	1188.500000	125.000000	360.00000	1.000000	1.000000
75%	5795.000000	2297.250000	164.750000	360.00000	1.000000	1.000000
max	81000.000000	41667.000000	700.000000	480.00000	1.000000	1.000000



```
In [8]: credit_df.Loan_Status.value_counts()
```


```
Out[8]: 1    422  
        0    192  
        Name: Loan_Status, dtype: int64
```

```
In [9]: credit_df.groupby(['Education', 'Loan_Status']).Education.count()
```

```
Out[9]: Education    Loan_Status  
Graduate           0         140  
                  1         340  
Not Graduate       0          52  
                  1          82  
        Name: Education, dtype: int64
```

```
In [10]: sns.barplot(y = 'Credit_History' , x='Loan_Status', hue='Education', data=credit_df)
```

```
Out[10]: <Axes: xlabel='Loan_Status', ylabel='Credit_History'>
```



```
In [11]: 100 * credit_df.isnull().sum() / credit_df.shape[0]
```

```
Out[11]: Loan_ID          0.000000
Gender          2.117264
Married         0.488599
Dependents      2.442997
Education       0.000000
Self_Employed   5.211726
ApplicantIncome 0.000000
CoapplicantIncome 0.000000
LoanAmount      0.000000
Loan_Amount_Term 2.280130
Credit_History  8.143322
Property_Area   0.000000
Loan_Status     0.000000
dtype: float64
```

```
In [12]: DF=credit_df.drop(credit_df.columns[0],axis=1)
```

```
In [13]: DF.head()
```

```
Out[13]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area
0	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	Urban
1	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	Rural
2	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	Urban
3	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	Urban
4	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	Urban

```
In [14]: object_columns = DF.select_dtypes(include=['object']).columns
numeric_columns = DF.select_dtypes(exclude=['object']).columns
```

```
In [15]: for column in object_columns:
majority = DF[column].value_counts().iloc[0]
DF[column].fillna(majority, inplace=True)
```

```
In [17]: for column in numeric_columns:
mean = DF[column].mean()
DF[column].fillna(mean, inplace=True)
```

```
In [18]: DF.head()
```

```
Out[18]:
```

	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area	Loan_Status
	No	0	Graduate	No	5849	0.0	0	360.0	1.0	Urban	1
	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	Rural	0
	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	Urban	1
	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	Urban	1
	No	0	Graduate	No	6000	0.0	141	360.0	1.0	Urban	1

```
In [20]: DF[object_columns].Property_Area #Categorical Columns
```

```
Out[20]: 0      Urban
          1      Rural
          2      Urban
          3      Urban
          4      Urban
          ...
          609    Rural
          610    Rural
          611    Urban
          612    Urban
          613    Semiurban
          Name: Property_Area, Length: 614, dtype: object
```

```
In [21]: DF[object_columns].Property_Area.head()
```

```
Out[21]: 0      Urban
          1      Rural
          2      Urban
          3      Urban
          4      Urban
          Name: Property_Area, dtype: object
```

```
In [22]: Df_dummy = pd.get_dummies(DF, columns=object_columns)
```

```
In [23]: Df_dummy.head()
```

```
Out[23]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status	Gender_489	Gender_Female	Gender_Male	Married_398
0	5849	0.0	0	360.0	1.0	1	0	0	1	0
1	4583	1508.0	128	360.0	1.0	0	0	0	1	0
2	3000	0.0	66	360.0	1.0	1	0	0	1	0
3	2583	2358.0	120	360.0	1.0	1	0	0	1	0
4	6000	0.0	141	360.0	1.0	1	0	0	1	0

5 rows × 25 columns

```
In [24]: Df_dummy.shape
```

```
Out[24]: (614, 25)
```

```
In [25]: from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LogisticRegression  
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [29]: x = Df_dummy.drop('Loan_Status', axis=1)  
y = Df_dummy.Loan_Status  
train_x, test_x, train_y, test_y = train_test_split(x, y, test_size=0.3, random_state=42)
```

```
In [30]: train_x.shape, test_x.shape
```

```
Out[30]: ((429, 24), (185, 24))
```

```
model = LogisticRegression() #LogisticRegression_Model
```

```
model.fit(train_x, train_y)
```

```
train_y_hat = model.predict(train_x)  
test_y_hat = model.predict(test_x)
```

```
print('train_accuracy', accuracy_score(train_y, train_y_hat))  
print('test accuracy', accuracy_score(test_y, test_y_hat))
```

```
train_accuracy 0.8205128205128205  
test accuracy 0.7837837837837838
```

```
In [35]: print(confusion_matrix(train_y, train_y_hat))
```

```
[[ 57  70]
 [   7 295]]
```

```
In [36]: print(confusion_matrix(test_y, test_y_hat))
```

```
[[ 27  38]
 [   2 118]]
```

```
In [37]: test_y.value_counts()
```

```
Out[37]: 1    120
         0     65
         Name: Loan_Status, dtype: int64
```

```
In [38]: pd.Series(test_y_hat).value_counts()
```

```
Out[38]: 1    156
         0     29
         dtype: int64
```

```
In [39]: (57 + 295) / train_y.shape[0]
```


```
Out[39]: 0.8205128205128205
```

```
In [40]: print(classification_report(test_y, test_y_hat))
```

	precision	recall	f1-score	support
0	0.93	0.42	0.57	65
1	0.76	0.98	0.86	120
accuracy			0.78	185
macro avg	0.84	0.70	0.71	185
weighted avg	0.82	0.78	0.76	185

```
In [41]: x = np.linspace(-10, 10, 100)
y = 1 / (1 + np.exp(-x)) #Sigmoid
plt.plot(x, y)
```

Out[41]: [<matplotlib.lines.Line2D at 0x1d457f4c460>]



## Practical-7

AIM: Program to Implement KNN Classification

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: credit_df = pd.read_csv('D:/Rukhsar_AIMLPracts/CreditRisk.csv')
```

```
In [3]: credit_df.shape
```

```
Out[3]: (614, 13)
```

```
In [4]: credit_df.head()
```

```
Out[4]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Pr
0	LP001002	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	

```
In [5]: credit_df.tail()
```

```
Out[5]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Pr
609	LP002978	Female	No	0	Graduate	No	2900	0.0	71	360.0	1.0	
610	LP002979	Male	Yes	3+	Graduate	No	4106	0.0	40	180.0	1.0	
611	LP002983	Male	Yes	1	Graduate	No	8072	240.0	253	360.0	1.0	
612	LP002984	Male	Yes	2	Graduate	No	7583	0.0	187	360.0	1.0	
613	LP002990	Female	No	0	Graduate	Yes	4583	0.0	133	360.0	1.0	

Activate Windows  
Go to Settings to activate

In [6]: credit\_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID               614 non-null    object
1   Gender                601 non-null    object
2   Married               611 non-null    object
3   Dependents            599 non-null    object
4   Education             614 non-null    object
5   Self_Employed         582 non-null    object
6   ApplicantIncome       614 non-null    int64
7   CoapplicantIncome     614 non-null    float64
8   LoanAmount            614 non-null    int64
9   Loan_Amount_Term      600 non-null    float64
10  Credit_History        564 non-null    float64
11  Property_Area         614 non-null    object
12  Loan_Status           614 non-null    int64
dtypes: float64(3), int64(3), object(7)
memory usage: 62.5+ KB
```

In [7]: credit\_df.describe()

Out[7]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status
count	614.000000	614.000000	614.000000	600.000000	564.000000	614.000000
mean	5403.459283	1621.245798	141.166124	342.000000	0.842199	0.687296
std	6109.041673	2926.248369	88.340630	65.12041	0.364878	0.463973
min	150.000000	0.000000	0.000000	12.000000	0.000000	0.000000
25%	2877.500000	0.000000	98.000000	360.000000	1.000000	0.000000
50%	3812.500000	1188.500000	125.000000	360.000000	1.000000	1.000000
75%	5795.000000	2297.250000	164.750000	360.000000	1.000000	1.000000
max	81000.000000	41667.000000	700.000000	480.000000	1.000000	1.000000

In [8]: credit\_df.Loan\_Status.value\_counts()

Out[8]: 1 422  
0 192  
Name: Loan\_Status, dtype: int64




```
In [9]: credit_df.groupby(['Education', 'Loan_Status']).Education.count()
```

```
Out[9]: Education    Loan_Status
Graduate          0          140
               1          340
Not Graduate      0           52
               1           82
Name: Education, dtype: int64
```

```
In [10]: sns.barplot(y = 'Credit_History', x='Loan_Status', hue='Education', data=credit_df)
```

```
Out[10]: <Axes: xlabel='Loan_Status', ylabel='Credit_History'>
```



```
In [11]: credit_df.isnull().sum()
```

```
Out[11]: Loan_ID          0
         Gender          13
         Married         3
         Dependents      15
         Education        0
         Self_Employed    32
         ApplicantIncome  0
         CoapplicantIncome 0
         LoanAmount       0
         Loan_Amount_Term 14
         Credit_History   50
         Property_Area    0
         Loan_Status      0
         dtype: int64
```

```
In [12]: DF=credit_df.drop(credit_df.columns[0],axis=1)
```

```
In [13]: DF.head()
```

```
Out[13]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area
0	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	Urban
1	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	Rural
2	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	Urban
3	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	Urban
4	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	Urban

```
In [14]: object_columns = DF.select_dtypes(include=['object']).columns
         numeric_columns = DF.select_dtypes(exclude=['object']).columns
```

```
In [15]: for column in object_columns:
         majority = DF[column].value_counts().iloc[0]
         DF[column].fillna(majority, inplace=True)
```

```
In [16]: for column in numeric_columns:
         mean = DF[column].mean()
         DF[column].fillna(mean, inplace=True)
```

```
In [17]: DF.head()
Out[17]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Property_Area
0	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	Urban
1	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	Rural
2	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	Urban
3	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	Urban
4	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	Urban

```
In [18]: DF[object_columns].Property_Area
Out[18]:
```

0	Urban
1	Rural
2	Urban
3	Urban
4	Urban
...	
609	Rural
610	Rural
611	Urban
612	Urban
613	Semiurban

Name: Property\_Area, Length: 614, dtype: object

Activate Windows  
Go to Settings to activate

```
In [19]: DF[object_columns].Property_Area.head()
Out[19]:
```

0	Urban
1	Rural
2	Urban
3	Urban
4	Urban

Name: Property\_Area, dtype: object

```
In [20]: Df_dummy = pd.get_dummies(DF, columns=object_columns)
```

```
In [21]: Df_dummy.head()
Out[21]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status	Gender_489	Gender_Female	Gender_Male	Married_398
0	5849	0.0	0	360.0	1.0	1	0	0	1	0
1	4583	1508.0	128	360.0	1.0	0	0	0	1	0
2	3000	0.0	66	360.0	1.0	1	0	0	1	0
3	2583	2358.0	120	360.0	1.0	1	0	0	1	0
4	6000	0.0	141	360.0	1.0	1	0	0	1	0

5 rows × 25 columns

Activate Windows  
Go to Settings to activate

In [22]: `Df_dummy.shape`

Out[22]: (614, 25)

In [31]: `from sklearn.model_selection import train_test_split  
from sklearn.neighbors import KNeighborsClassifier  
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report`

In [24]: `x = Df_dummy.drop('Loan_Status', axis=1)  
y = Df_dummy.Loan_Status  
train_x, test_x, train_y, test_y = train_test_split(x, y, test_size=0.3, random_state=42)`

In [25]: `train_x.shape, test_x.shape`

Out[25]: ((429, 24), (185, 24))

In [26]: `knn_model = KNeighborsClassifier(n_neighbors=7) #KNN_Model`

In [27]: `knn_model.fit(train_x, train_y)`

Out[27]: 

▼ KNeighborsClassifier

KNeighborsClassifier(n\_neighbors=7)

```
In [32]: train_y_hat = knn_model.predict(train_x)
test_y_hat = knn_model.predict(test_x)

print('-' * 20, 'Train', '-' * 20)
print(classification_report(train_y, train_y_hat))
print('-' * 20, 'Test', '-' * 20)
print(classification_report(test_y, test_y_hat))
```

```
----- Train -----
              precision    recall  f1-score   support

         0           0.70       0.24       0.35         127
         1           0.75       0.96       0.84         302

   accuracy                   0.74         429
  macro avg           0.72       0.60       0.60         429
 weighted avg           0.73       0.74       0.70         429

----- Test -----
              precision    recall  f1-score   support

         0           0.36       0.12       0.18          65
         1           0.65       0.88       0.75         120

   accuracy                   0.62         185
  macro avg           0.51       0.50       0.47         185
 weighted avg           0.55       0.62       0.55         185
```

## Practical-8

AIM: Program to Implement Principal Component Analysis

# Principal Component Analysis

```
In [1]: from sklearn.datasets import load_digits
import pandas as pd
dataset=load_digits()
dataset.keys()
```

```
Out[1]: dict_keys(['data', 'target', 'frame', 'feature_names', 'target_names', 'images', 'DESCR'])
```

```
In [2]: dataset.data.shape
```

```
Out[2]: (1797, 64)
```

```
In [3]: dataset.data[0]
```

```
Out[3]: array([ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.,  0.,  0., 13., 15., 10.,
                15.,  5.,  0.,  0.,  3., 15.,  2.,  0., 11.,  8.,  0.,  0.,  4.,
                12.,  0.,  0.,  8.,  8.,  0.,  0.,  5.,  8.,  0.,  0.,  9.,  8.,
                 0.,  0.,  4., 11.,  0.,  1., 12.,  7.,  0.,  0.,  2., 14.,  5.,
                10., 12.,  0.,  0.,  0.,  0.,  6., 13., 10.,  0.,  0.,  0.])
```


```
In [4]: dataset.data[0].reshape(8,8)
```

```
Out[4]: array([[ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.],
                [ 0.,  0., 13., 15., 10., 15.,  5.,  0.],
                [ 0.,  3., 15.,  2.,  0., 11.,  8.,  0.],
                [ 0.,  4., 12.,  0.,  0.,  8.,  8.,  0.],
                [ 0.,  5.,  8.,  0.,  0.,  9.,  8.,  0.],
                [ 0.,  4., 11.,  0.,  1., 12.,  7.,  0.],
                [ 0.,  2., 14.,  5., 10., 12.,  0.,  0.],
                [ 0.,  0.,  6., 13., 10.,  0.,  0.,  0.]])
```

```
In [5]: from matplotlib import pyplot as plt
import matplotlib inline
plt.gray()
plt.matshow(dataset.data[0].reshape(8,8))
```


```
Out[5]: <matplotlib.image.AxesImage at 0x1c9c48eeda0>

<Figure size 640x480 with 0 Axes>
```



```
In [6]: plt.matshow(dataset.data[9].reshape(8,8))
```

```
Out[6]: <matplotlib.image.AxesImage at 0x1c9c4a21900>
```



```
In [7]: dataset.target[:5]
```

```
Out[7]: array([0, 1, 2, 3, 4])
```

```
In [8]: df=pd.DataFrame(dataset.data,columns=dataset.feature_names)
df.head()
```

```
Out[8]:
```

	pixel_0_0	pixel_0_1	pixel_0_2	pixel_0_3	pixel_0_4	pixel_0_5	pixel_0_6	pixel_0_7	pixel_1_0	pixel_1_1	...	pixel_6_6	pixel_6_7	pixel_7_0	pixel_7_1	pix
0	0.0	0.0	5.0	13.0	9.0	1.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	
1	0.0	0.0	0.0	12.0	13.0	5.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	
2	0.0	0.0	0.0	4.0	15.0	12.0	0.0	0.0	0.0	0.0	...	5.0	0.0	0.0	0.0	
3	0.0	0.0	7.0	15.0	13.0	1.0	0.0	0.0	0.0	8.0	...	9.0	0.0	0.0	0.0	
4	0.0	0.0	0.0	1.0	11.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	

5 rows x 64 columns

```
In [9]: dataset.target
```

```
Out[9]: array([0, 1, 2, ..., 8, 9, 8])
```

```
In [10]: df.describe()
```

```
Out[10]:
```

	pixel_0_0	pixel_0_1	pixel_0_2	pixel_0_3	pixel_0_4	pixel_0_5	pixel_0_6	pixel_0_7	pixel_1_0	pixel_1_1	...	pixel_6_6	pix
count	1797.0	1797.000000	1797.000000	1797.000000	1797.000000	1797.000000	1797.000000	1797.000000	1797.000000	1797.000000	...	1797.000000	1797.
mean	0.0	0.303840	5.204786	11.835838	11.848080	5.781859	1.362270	0.129661	0.005565	1.993879	...	3.725097	0.
std	0.0	0.907192	4.754826	4.248842	4.287388	5.666418	3.325775	1.037383	0.094222	3.196160	...	4.919406	0.
min	0.0	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	...	0.000000	0.
25%	0.0	0.000000	1.000000	10.000000	10.000000	0.000000	0.000000	0.000000	0.000000	0.000000	...	0.000000	0.
50%	0.0	0.000000	4.000000	13.000000	13.000000	4.000000	0.000000	0.000000	0.000000	0.000000	...	1.000000	0.
75%	0.0	0.000000	9.000000	15.000000	15.000000	11.000000	0.000000	0.000000	0.000000	0.000000	...	7.000000	0.
max	0.0	8.000000	16.000000	16.000000	16.000000	16.000000	16.000000	15.000000	2.000000	16.000000	...	16.000000	13.

8 rows x 64 columns

```
In [11]: x=df
y=dataset.target
```

```
In [12]: from sklearn.preprocessing import StandardScaler
```

```
In [13]: scaler=StandardScaler()
x_scaled=scaler.fit_transform(x)
x_scaled
```

```
Out[13]: array([[ 0.          , -0.33501649, -0.04308102, ..., -1.14664746,
-0.5056698 , -0.19600752],
[ 0.          , -0.33501649, -1.09493684, ...,  0.54856067,
-0.5056698 , -0.19600752],
[ 0.          , -0.33501649, -1.09493684, ...,  1.56568555,
 1.6951369 , -0.19600752],
...,
[ 0.          , -0.33501649, -0.88456568, ..., -0.12952258,
-0.5056698 , -0.19600752],
[ 0.          , -0.33501649, -0.67419451, ...,  0.8876023 ,
-0.5056698 , -0.19600752],
[ 0.          , -0.33501649, 1.00877481, ...,  0.8876023 ,
-0.26113572, -0.19600752]])
```

```
In [14]: from sklearn.model_selection import train_test_split
```

```
In [16]: x_train,x_test,y_train,y_test=train_test_split(x_scaled,y,test_size=0.2, random_state=30)
```

```
In [17]: from sklearn.linear_model import LogisticRegression
```

```
In [19]: model = LogisticRegression()
model.fit(x_train, y_train)
model.score(x_test, y_test)
```



Out[19]: 0.9722222222222222

In [20]: `from sklearn.decomposition import PCA`

In [21]: `pca = PCA(0.95)`  
`x_pca = pca.fit_transform(x)`  
`x_pca.shape`

Out[21]: (1797, 29)

In [22]: `pca.explained_variance_ratio_`

Out[22]: array([[0.14890594, 0.13618771, 0.11794594, 0.08409979, 0.05782415,  
0.0491691 , 0.04315987, 0.03661373, 0.03353248, 0.03078806,  
0.02372341, 0.02272697, 0.01821863, 0.01773855, 0.01467101,  
0.01409716, 0.01318589, 0.01248138, 0.01017718, 0.00905617,  
0.00889538, 0.00797123, 0.00767493, 0.00722904, 0.00695889,  
0.00596081, 0.00575615, 0.00515158, 0.0048954 ])

In [23]: `pca.n_components_`

Out[23]: 29

In [24]: `x_train_pca, x_test_pca, y_train, y_test = train_test_split(x_pca, y, test_size=0.2, random_state=30)`

In [25]: `from sklearn.linear_model import LogisticRegression`

In [26]: `model = LogisticRegression(max_iter=1000)`  
`model.fit(x_train_pca, y_train)`  
`model.score(x_test_pca, y_test)`

C:\ProgramData\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
n\_iter\_i = \_check\_optimize\_result(

Out[26]: 0.9694444444444444

In [28]: `pca = PCA(n_components=2)`  
`x_pca = pca.fit_transform(x)`  
`x_pca.shape`

Out[28]: (1797, 2)

In [29]: `pca.explained_variance_ratio_`

Out[29]: array([0.14890594, 0.13618771])

In [30]: `x_train_pca, x_test_pca, y_train, y_test = train_test_split(x_pca, y, test_size=0.2, random_state=30)`

In [31]: `model = LogisticRegression(max_iter=1000)`  
`model.fit(x_train_pca, y_train)`  
`model.score(x_test_pca, y_test)`

Out[31]: 0.6083333333333333

## Practical-9

AIM: Program to Implement K means Algorithm

```
In [1]: import pandas as pd
        from sklearn.cluster import KMeans
        import matplotlib.pyplot as plt
```

```
In [3]: df = pd.read_csv('D:/Rukhsar_AIMLPracts/dataset2.csv')
        df.head()
```

```
Out[3]:
```

	id	mean_dist_day	mean_over_speed_perc
0	3423311935	71.24	28
1	3423313212	52.53	25
2	3423313724	64.54	27
3	3423311373	55.69	22
4	3423310999	54.58	25

```
In [4]: df.describe()
```

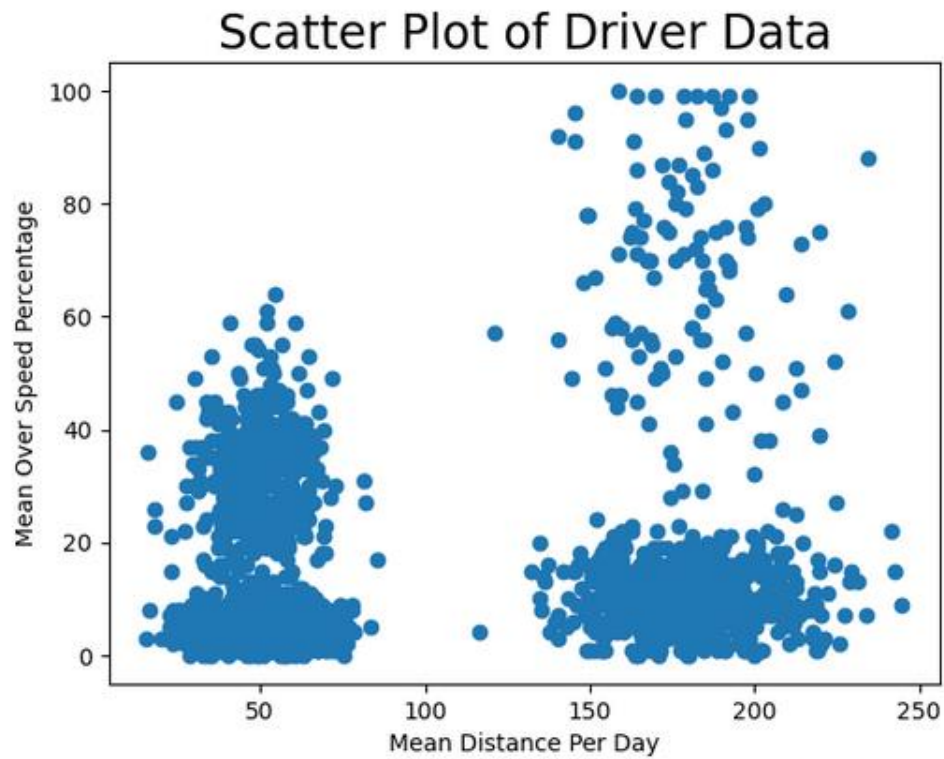
```
Out[4]:
```

	id	mean_dist_day	mean_over_speed_perc
count	4.000000e+03	4000.000000	4000.000000
mean	3.423312e+09	76.041522	10.721000
std	1.154845e+03	53.469563	13.708543
min	3.423310e+09	15.520000	0.000000
25%	3.423311e+09	45.247500	4.000000
50%	3.423312e+09	53.330000	6.000000
75%	3.423313e+09	65.632500	9.000000
max	3.423314e+09	244.790000	100.000000

```
In [5]: df.shape
```

```
Out[5]: (4000, 3)
```

```
In [6]: plt.plot(df.mean_dist_day, df.mean_over_speed_perc, 'o')  
plt.xlabel('Mean Distance Per Day')  
plt.ylabel('Mean Over Speed Percentage')  
plt.title('Scatter Plot of Driver Data', fontsize=20)  
plt.show()
```



In [7]: df.head()

Out[7]:

	id	mean_dist_day	mean_over_speed_perc
0	3423311935	71.24	28
1	3423313212	52.53	25
2	3423313724	64.54	27
3	3423311373	55.69	22
4	3423310999	54.58	25

In [8]: data = df.drop(['id'], axis=1)  
cluster\_model = KMeans(n\_clusters=2)  
cluster\_model.fit(data)

Out[8]:

KMeans

KMeans(n\_clusters=2)

In [9]: df['labels'] = cluster\_model.labels\_

In [10]: df.head()

Out[10]:

	id	mean_dist_day	mean_over_speed_perc	labels
0	3423311935	71.24	28	1
1	3423313212	52.53	25	1
2	3423313724	64.54	27	1
3	3423311373	55.69	22	1
4	3423310999	54.58	25	1


In [11]: df.labels.unique()

Out[11]: array([1, 0])

In [12]: df['labels'].value\_counts()

Out[12]: 1     3200  
          0     800  
          Name: labels, dtype: int64


```
In [18]: for label in df.labels.unique():  
         plt.plot(df.loc[df.labels == label, 'mean_dist_day'],  
                  df.loc[df.labels == label, 'mean_over_speed_perc'], 'o', label=label)  
plt.xlabel('Mean Distance Per Day')  
plt.ylabel('Mean Over Speed Percentage')  
plt.title('Scatter Plot of Drivers Data', fontsize=20)  
plt.legend()  
plt.show()
```



```
In [19]: cluster_model.cluster_centers_
```

```
Out[19]: array([[180.017075 , 18.29      ],  
               [ 50.04763438,  8.82875   ]])
```

```
In [20]: error = []  
for k in range(1,11):  
    cluster_model = KMeans(k)  
    cluster_model.fit(data)  
    error.append(cluster_model.inertia_)  
plt.plot(range(1,11), error)  
plt.show()
```



## Practical-10

AIM: Program to implement Support Vector Machine (SVM)

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: credit_df = pd.read_csv('D:/Rukhsar_AIMLPracts/CreditRisk.csv')
```

```
In [3]: credit_df.shape
```

```
Out[3]: (614, 13)
```

```
In [4]: credit_df.head()
```

```
Out[4]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_Score
0	LP001002	Male	No	0	Graduate	No	5849	0.0	0	360.0	
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141	360.0	

```
In [5]: credit_df.tail()
```

```
Out[5]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_Score
609	LP002978	Female	No	0	Graduate	No	2900	0.0	71	360.0	
610	LP002979	Male	Yes	3+	Graduate	No	4106	0.0	40	180.0	
611	LP002983	Male	Yes	1	Graduate	No	8072	240.0	253	360.0	
612	LP002984	Male	Yes	2	Graduate	No	7583	0.0	187	360.0	
613	LP002990	Female	No	0	Graduate	Yes	4583	0.0	133	360.0	

In [6]: credit\_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID                614 non-null    object
1   Gender                 601 non-null    object
2   Married                611 non-null    object
3   Dependents             599 non-null    object
4   Education              614 non-null    object
5   Self_Employed          582 non-null    object
6   ApplicantIncome        614 non-null    int64
7   CoapplicantIncome      614 non-null    float64
8   LoanAmount             614 non-null    int64
9   Loan_Amount_Term       600 non-null    float64
10  Credit_History         564 non-null    float64
11  Property_Area          614 non-null    object
12  Loan_Status            614 non-null    int64
dtypes: float64(3), int64(3), object(7)
memory usage: 62.5+ KB
```

In [7]: credit\_df.describe()

Out[7]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status
count	614.000000	614.000000	614.000000	600.00000	564.000000	614.000000
mean	5403.459283	1621.245798	141.166124	342.00000	0.842199	0.687296
std	6109.041673	2926.248369	88.340630	65.12041	0.364878	0.463973
min	150.000000	0.000000	0.000000	12.00000	0.000000	0.000000
25%	2877.500000	0.000000	98.000000	360.00000	1.000000	0.000000
50%	3812.500000	1188.500000	125.000000	360.00000	1.000000	1.000000
75%	5795.000000	2297.250000	164.750000	360.00000	1.000000	1.000000
max	81000.000000	41667.000000	700.000000	480.00000	1.000000	1.000000

In [8]: credit\_df.Loan\_Status.value\_counts()

Out[8]:

```
1    422
0    192
Name: Loan_Status, dtype: int64
```




```
In [9]: credit_df.groupby(['Education', 'Loan_Status']).Education.count()
```

```
Out[9]: Education    Loan_Status
Graduate          0          140
              1          340
Not Graduate      0           52
              1           82
Name: Education, dtype: int64
```

```
In [10]: sns.barplot(y = 'Credit_History' , x='Loan_Status', hue='Education', data=credit_df)
```

```
Out[10]: <Axes: xlabel='Loan_Status', ylabel='Credit_History'>
```



In [11]: credit\_df.isnull().sum()

Out[11]:

Loan_ID	0
Gender	13
Married	3
Dependents	15
Education	0
Self_Employed	32
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	0
Loan_Amount_Term	14
Credit_History	50
Property_Area	0
Loan_Status	0
dtype:	int64

In [13]: DF.head()

Out[13]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Pr
0	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	
1	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	
2	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	
3	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	
4	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	

In [14]:

```
object_columns = DF.select_dtypes(include=['object']).columns
numeric_columns = DF.select_dtypes(exclude=['object']).columns
```

In [15]:

```
for column in object_columns:
    majority = DF[column].value_counts().iloc[0]
    DF[column].fillna(majority, inplace=True)
```

In [16]:

```
for column in numeric_columns:
    mean = DF[column].mean()
    DF[column].fillna(mean, inplace=True)
```

In [17]: DF.head()

Out[17]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Pr
0	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	
1	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	
2	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	
3	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	
4	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	

```
In [18]: credit_df.drop('Loan_ID', axis=1, inplace=True)
```

```
object_columns=credit_df.select_dtypes(include=['object']).columns
```

```
In [19]: object_columns=credit_df.select_dtypes(include=['object']).columns
```

```
In [20]: credit_df.head()
```

```
Out[20]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Pr
0	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	
1	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	
2	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	
3	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	
4	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	

```
In [21]: DF[object_columns].Property_Area
```

```
Out[21]:
```

0	Urban
1	Rural
2	Urban
3	Urban
4	Urban
...	
609	Rural
610	Rural
611	Urban
612	Urban
613	Semiurban

Name: Property\_Area, Length: 614, dtype: object

```
In [22]: DF[object_columns].Property_Area.head()
```

```
Out[22]:
```

0	Urban
1	Rural
2	Urban
3	Urban
4	Urban

Name: Property\_Area, dtype: object

```
In [23]: Df_dummy = pd.get_dummies(DF, columns=object_columns)
```

In [24]: Df\_dummy.head()

Out[24]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status	Gender_Male	Gender_Female	Gender_489	Ma
0	5849	0.0	0	360.0	1.0	1	0	0	0	1
1	4583	1508.0	128	360.0	1.0	0	0	0	0	1
2	3000	0.0	66	360.0	1.0	1	0	0	0	1
3	2583	2358.0	120	360.0	1.0	1	0	0	0	1
4	6000	0.0	141	360.0	1.0	1	0	0	0	1

5 rows × 25 columns

< >

In [25]: Df\_dummy.shape

Out[25]: (614, 25)

In [27]: `from sklearn.model_selection import train_test_split`  
`from sklearn.svm import SVC`  
`from sklearn.metrics import accuracy_score, confusion_matrix, classification_report`

In [30]: `x = Df_dummy.drop('Loan_Status', axis=1)`  
`y = Df_dummy.Loan_Status`  
`train_x, test_x, train_y, test_y = train_test_split(x, y, test_size=0.3, random_state=42)`

In [31]: `train_x.shape, test_x.shape`

Out[31]: ((429, 24), (185, 24))

Activate Win  
Go to Settings to

In [27]: `from sklearn.model_selection import train_test_split`  
`from sklearn.svm import SVC`  
`from sklearn.metrics import accuracy_score, confusion_matrix, classification_report`

In [30]: `x = Df_dummy.drop('Loan_Status', axis=1)`  
`y = Df_dummy.Loan_Status`  
`train_x, test_x, train_y, test_y = train_test_split(x, y, test_size=0.3, random_state=42)`

In [31]: `train_x.shape, test_x.shape`

Out[31]: ((429, 24), (185, 24))

In [32]: `svm_model = SVC(kernel='rbf', gamma=0.00001, C=1000)`

In [33]: `svm_model.fit(train_x, train_y)`

```
In [34]: train_y_hat = svm_model.predict(train_x)
test_y_hat = svm_model.predict(test_x)
```

```
In [35]: print('-'*20, 'Train', '-'*20)
print(classification_report(train_y, train_y_hat))
print('-'*20, 'Test', '-'*20)
print(classification_report(test_y, test_y_hat))
```

```
----- Train -----
              precision    recall  f1-score   support

      0       0.95       0.95       0.95        127
      1       0.98       0.98       0.98        302

   accuracy                0.97        429
  macro avg       0.96       0.96       0.96        429
 weighted avg       0.97       0.97       0.97        429

----- Test -----
              precision    recall  f1-score   support

      0       0.36       0.18       0.24         65
      1       0.65       0.82       0.73        120

   accuracy                0.60        185
  macro avg       0.51       0.50       0.49        185
 weighted avg       0.55       0.60       0.56        185
```

```
In [36]: confusion_matrix(test_y, test_y_hat)
```

```
Out[36]: array([[12, 53],
                [21, 99]], dtype=int64)
```

# Practical 11

AIM: Program to Implement Decision Tree Algorithm

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: credit_df = pd.read_csv('D:/Rukhsar_AIMLPPracts/CreditRisk.csv')
```

```
In [3]: credit_df.shape
```

```
Out[3]: (614, 13)
```

```
In [4]: credit_df.head()
```

```
Out[4]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_L
0	LP001002	Male	No	0	Graduate	No	5849	0.0	0	360.0	
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141	360.0	

```
In [5]: credit_df.tail()
```

```
Out[5]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_L
609	LP002978	Female	No	0	Graduate	No	2900	0.0	71	360.0	
610	LP002979	Male	Yes	3+	Graduate	No	4106	0.0	40	180.0	
611	LP002983	Male	Yes	1	Graduate	No	8072	240.0	253	360.0	
612	LP002984	Male	Yes	2	Graduate	No	7583	0.0	187	360.0	
613	LP002990	Female	No	0	Graduate	Yes	4583	0.0	133	360.0	

In [6]: credit\_df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID               614 non-null    object
1   Gender                601 non-null    object
2   Married               611 non-null    object
3   Dependents            599 non-null    object
4   Education             614 non-null    object
5   Self_Employed         582 non-null    object
6   ApplicantIncome       614 non-null    int64
7   CoapplicantIncome     614 non-null    float64
8   LoanAmount            614 non-null    int64
9   Loan_Amount_Term      600 non-null    float64
10  Credit_History        564 non-null    float64
11  Property_Area         614 non-null    object
12  Loan_Status           614 non-null    int64
dtypes: float64(3), int64(3), object(7)
memory usage: 62.5+ KB
```

In [7]: credit\_df.describe()

Out[7]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status
count	614.000000	614.000000	614.000000	600.00000	564.000000	614.000000
mean	5403.459283	1621.245798	141.166124	342.00000	0.842199	0.687296
std	6109.041673	2926.248369	88.340630	65.12041	0.364878	0.463973
min	150.000000	0.000000	0.000000	12.00000	0.000000	0.000000
25%	2877.500000	0.000000	98.000000	360.00000	1.000000	0.000000
50%	3812.500000	1188.500000	125.000000	360.00000	1.000000	1.000000
75%	5795.000000	2297.250000	164.750000	360.00000	1.000000	1.000000
max	81000.000000	41667.000000	700.000000	480.00000	1.000000	1.000000

In [8]: credit\_df.Loan\_Status.value\_counts()

Out[8]:


```
1    422
0    192
Name: Loan_Status, dtype: int64
```

```
In [9]: credit_df.groupby(['Education', 'Loan_Status']).Education.count()
```

```
Out[9]: Education    Loan_Status
Graduate           0         140
                1         340
Not Graduate       0          52
                1          82
Name: Education, dtype: int64
```

```
In [10]: sns.barplot(y = 'Credit_History', x='Loan_Status', hue='Education', data=credit_df)
```

```
Out[10]: <Axes: xlabel='Loan_Status', ylabel='Credit_History'>
```





```
In [11]: credit_df.isnull().sum()
```

```
Out[11]: Loan_ID          0
Gender          13
Married         3
Dependents      15
Education       0
Self_Employed   32
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount      0
Loan_Amount_Term 14
Credit_History  50
Property_Area   0
Loan_Status     0
dtype: int64
```

```
In [12]: DF=credit_df.drop(credit_df.columns[0],axis=1)
```

```
In [13]: DF.head()
```

```
Out[13]:
```

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Pr
0	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	
1	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	
2	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	
3	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	
4	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	

```
In [14]: object_columns = DF.select_dtypes(include=['object']).columns
numeric_columns = DF.select_dtypes(exclude=['object']).columns
```

```
In [15]: for column in object_columns:
majority = DF[column].value_counts().iloc[0]
DF[column].fillna(majority, inplace=True)
```

```
In [16]: for column in numeric_columns:
mean = DF[column].mean()
DF[column].fillna(mean, inplace=True)
```

In [17]: DF.head()

Out[17]:

	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Pr
0	Male	No	0	Graduate	No	5849	0.0	0	360.0	1.0	
1	Male	Yes	1	Graduate	No	4583	1508.0	128	360.0	1.0	
2	Male	Yes	0	Graduate	Yes	3000	0.0	66	360.0	1.0	
3	Male	Yes	0	Not Graduate	No	2583	2358.0	120	360.0	1.0	
4	Male	No	0	Graduate	No	6000	0.0	141	360.0	1.0	

In [18]: DF[object\_columns].Property\_Area

Out[18]:

```
0      Urban
1      Rural
2      Urban
3      Urban
4      Urban
...
609     Rural
610     Rural
611     Urban
612     Urban
613  Semiurban
Name: Property_Area, Length: 614, dtype: object
```

Activate Wir

In [19]: DF[object\_columns].Property\_Area.head()

Out[19]:

```
0      Urban
1      Rural
2      Urban
3      Urban
4      Urban
Name: Property_Area, dtype: object
```

In [20]: Df\_dummy = pd.get\_dummies(DF, columns=object\_columns)

In [21]: Df\_dummy.head()

Out[21]:

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Loan_Status	Gender_489	Gender_Female	Gender_Male	Ma
0	5849	0.0	0	360.0	1.0	1	0	0	1	
1	4583	1508.0	128	360.0	1.0	0	0	0	1	
2	3000	0.0	66	360.0	1.0	1	0	0	1	
3	2583	2358.0	120	360.0	1.0	1	0	0	1	
4	6000	0.0	141	360.0	1.0	1	0	0	1	

5 rows × 25 columns

```
In [16]: from sklearn.model_selection import train_test_split as TTS
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
```

```
In [17]: X = DF_dummy.drop('Loan_Status', axis=1)
Y = DF_dummy.Loan_Status
train_x, test_x, train_y, test_y = TTS(X, Y, test_size = 0.3, random_state=42)
```

```
In [18]: train_x.shape, test_x.shape
```

```
Out[18]: ((429, 24), (185, 24))
```

```
In [19]: from sklearn.tree import DecisionTreeClassifier
dt_model = DecisionTreeClassifier(max_depth=14)
```

```
In [20]: dt_model.fit(train_x, train_y)
```

```
Out[20]: DecisionTreeClassifier
DecisionTreeClassifier(max_depth=14)
```

```
In [21]: train_y_hat = dt_model.predict(train_x)
test_y_hat = dt_model.predict(test_x)
```

```
In [22]: print('-'*20, 'Train', '-'*20)
print(classification_report(train_y, train_y_hat))
print('-'*20, 'Test', '-'*20)
print(classification_report(test_y, test_y_hat))
```

```
----- Train -----
              precision    recall  f1-score   support

      0      0.99      1.00      1.00       127
      1      1.00      1.00      1.00       302

   accuracy                1.00       429
  macro avg              1.00      1.00      1.00       429
 weighted avg              1.00      1.00      1.00       429

----- Test -----
              precision    recall  f1-score   support

      0      0.57      0.54      0.56        65
      1      0.76      0.78      0.77       120

   accuracy                0.70       185
  macro avg              0.67      0.66      0.66       185
 weighted avg              0.69      0.70      0.69       185
```

```
In [23]: confusion_matrix(train_y, train_y_hat)
```

```
Out[23]: array([[127,  0],
                [ 1, 301]], dtype=int64)
```

## Practical 12

AIM: Program to Implement Random Forest

```
In [1]: import numpy as np
import pandas as pd
import matplotlib as mpl
import matplotlib.pyplot as plt
```

```
In [2]: train = pd.read_csv(r"F:\FYMCA\Sem2\AIML\titanic.csv")
print(train.shape)

(891, 12)
```

```
In [3]: #checking for missing data
NAs = pd.concat([train.isnull().sum()], axis=1, keys=["Train"])
NAs[NAs.sum(axis=1) > 0]
```

Out[3]:

	Train
Age	177
Cabin	687
Embarked	2

```
In [4]: train.pop("Cabin")
train.pop("Name")
train.pop("Ticket")
```

```
Out[4]: 0      A/5  21171
1      PC  17599
2  STON/O2. 3101282
3      113803
4      373450
...
886      211536
887      112053
888  W./C.  6607
889      111369
890      370376
Name: Ticket, Length: 891, dtype: object
```

```
In [5]: # Filling missing Age values with mean
train["Age"] = train["Age"].fillna(train["Age"].mean())
```

```
In [6]: # Filling missing Embarked values with most common value
train["Embarked"] = train["Embarked"].fillna(train["Embarked"].mode()[0])
```

```
In [7]: train["Pclass"] = train["Pclass"].apply(str)
```

```
In [8]: # Getting Dummies from all other categorical vars
for col in train.dtypes[train.dtypes == "object"].index:
    for_dummy = train.pop(col)
    train = pd.concat([train, pd.get_dummies(for_dummy, prefix=col)], axis=1)
train.head()
```

```
Out[8]:
```

	PassengerId	Survived	Age	SibSp	Parch	Fare	Pclass_1	Pclass_2	Pclass_3	Sex_female	Sex_male	Embarked_C	Embarked_Q	Embarked_S
0	1	0	22.0	1	0	7.2500	0	0	1	0	1	0	0	1
1	2	1	38.0	1	0	71.2833	1	0	0	1	0	1	0	0
2	3	1	26.0	0	0	7.9250	0	0	1	1	0	0	0	1
3	4	1	35.0	1	0	53.1000	1	0	0	1	0	0	0	1
4	5	0	35.0	0	0	8.0500	0	0	1	0	1	0	0	1

```
In [9]: labels = train.pop("Survived")
```

```
In [10]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(train, labels, test_size=0.25)
```

```
In [11]: from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(n_estimators=100)
rf.fit(x_train, y_train)
```

```
Out[11]: RandomForestClassifier
RandomForestClassifier()
```

```
In [12]: y_pred = rf.predict(x_test)
```


```
In [14]: from sklearn.metrics import roc_curve, auc
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(false_positive_rate, true_positive_rate)
roc_auc
```

```
Out[14]: 0.7425986842105263
```

```
In [15]: n_estimators = [1, 2, 4, 8, 16, 32, 64, 100, 200]
train_results = []
test_results = []
```

```
In [16]: for estimator in n_estimators:
    rf = RandomForestClassifier(n_estimators=estimator, n_jobs=-1)
    rf.fit(x_train, y_train)
    train_pred = rf.predict(x_train)
    false_positive_rate, true_positive_rate, thresholds = roc_curve(y_train, train_pred)
    roc_auc = auc(false_positive_rate, true_positive_rate)
    train_results.append(roc_auc)
    y_pred = rf.predict(x_test)
    false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred)
    roc_auc = auc(false_positive_rate, true_positive_rate)
    test_results.append(roc_auc)
```

```
In [17]: from matplotlib.legend_handler import HandlerLine2D
line1, = plt.plot(n_estimators, train_results, "b", label="Train AUC")
line2, = plt.plot(n_estimators, test_results, "r", label="Test AUC")
plt.legend(handler_map={line1: HandlerLine2D(numpoints=2)})
plt.ylabel("AUC score")
plt.xlabel("n_estimators")
plt.show()
```



```
In [18]: from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(n_estimators=200)
rf.fit(x_train, y_train)
```

```
Out[18]: RandomForestClassifier
RandomForestClassifier(n_estimators=200)
```


```
In [19]: y_pred = rf.predict(x_test)
from sklearn.metrics import roc_curve, auc
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(false_positive_rate, true_positive_rate)
roc_auc
```

```
Out[19]: 0.7400493421052632
```

```
In [20]: n_estimators = [1, 2, 4, 8, 16, 32, 64, 100, 200]
train_results = []
test_results = []
```

```
In [21]: for estimator in n_estimators:
rf = RandomForestClassifier(n_estimators=estimator, n_jobs=-1)
rf.fit(x_train, y_train)
train_pred = rf.predict(x_train)
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_train, train_pred)
roc_auc = auc(false_positive_rate, true_positive_rate)
train_results.append(roc_auc)
y_pred = rf.predict(x_test)
false_positive_rate, true_positive_rate, thresholds = roc_curve(y_test, y_pred)
roc_auc = auc(false_positive_rate, true_positive_rate)
test_results.append(roc_auc)
```

```
In [22]: from matplotlib.legend_handler import HandlerLine2D
line1, = plt.plot(n_estimators, train_results, "b", label="Train AUC")
line2, = plt.plot(n_estimators, test_results, "r", label="Test AUC")
plt.legend(handler_map={line1: HandlerLine2D(numpoints=2)})
plt.ylabel("AUC score")
plt.xlabel("n_estimators")
plt.show()
```



## Practical 13

AIM: Program to Implement AdaBoost

```
[1]: from sklearn.ensemble import AdaBoostClassifier
[2]: from sklearn import datasets
[3]: from sklearn.model_selection import train_test_split
[4]: from sklearn import metrics
[5]: iris=datasets.load_iris()
    X=iris.data
    y=iris.target
[6]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2)
[7]: AdaModel=AdaBoostClassifier(n_estimators=100,learning_rate=1)
    model=AdaModel.fit(X_train,y_train)
    y_pred=model.predict(X_test)
[11]: print("Accuracy:",metrics.accuracy_score(y_test,y_pred))

Accuracy: 0.9333333333333333
```

```
[11]: print("Accuracy:",metrics.accuracy_score(y_test,y_pred))

Accuracy: 0.9333333333333333
```

```
[25]: from sklearn.svm import SVC
    from sklearn import metrics
    svc=SVC(probability=True,kernel='linear')
    abc=AdaBoostClassifier(n_estimators=70,base_estimator=svc,learning_rate=1)
```

```
[26]: model=abc.fit(X_train,y_train)
    y_pred=model.predict(X_test)
```

```
C:\Users\USER\anaconda3\Lib\site-packages\sklearn\ensemble\_base.py:166: FutureWarning: `base_estimator` was renamed to `estimator` in version 1.2 and will be removed in 1.4.
  warnings.warn(
```

```
[27]: print("Accuracy:",metrics.accuracy_score(y_test,y_pred))

Accuracy: 0.9
```



# Practical 14

AIM: Program to Implement Gradient Boosting

```
In [2]: # Importing necessary packages
from sklearn.ensemble import GradientBoostingRegressor
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import r2_score
import warnings

# Suppressing warnings
warnings.filterwarnings('ignore')
```

```
In [4]: # Fetching the California housing dataset
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()

# Splitting the data into features and target variable
X = pd.DataFrame(housing.data, columns=housing.feature_names)
y = pd.DataFrame(housing.target, columns=['target'])
```

```
In [5]: #Viewing Data - predictors
X.head()
```

```
Out[5]:
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23

```
In [5]: #Viewing Data - predictors
X.head()
```

```
Out[5]:
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup	Latitude	Longitude
0	8.3252	41.0	6.984127	1.023810	322.0	2.555556	37.88	-122.23
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842	37.86	-122.22
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260	37.85	-122.24
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945	37.85	-122.25
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467	37.85	-122.25

```
In [ ]: y[1:10] #response
```

```
In [7]: # Split dataset into training set and test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2) # 80% training and 20% test
```

```
In [8]: # Create gradientboost REGRESSOR object
gradientregressor = GradientBoostingRegressor(max_depth=2,n_estimators=3,learning_rate=1.0)
```

```
In [9]: # Train gradientboost REGRESSOR
model = gradientregressor.fit(X_train, y_train)


#Predict the response for test dataset
y_pred = model.predict(X_test)
```

```
In [ ]: r2_score(y_pred,y_test)|
```

```
In [14]: import matplotlib.pyplot as plt
%matplotlib inline

# Plot feature importance
feature_importance = model.feature_importances_

# make importances relative to max importance
feature_importance = 100.0 * (feature_importance / feature_importance.max())
sorted_idx = np.argsort(feature_importance)
pos = np.arange(sorted_idx.shape[0]) + .5
plt.barh(pos, feature_importance[sorted_idx], align='center')
plt.yticks(pos, boston.feature_names[sorted_idx])
plt.xlabel('Relative Importance')
plt.title('Variable Importance')
plt.show()
```



```
In [*]: from sklearn.model_selection import GridSearchCV
LR = {'learning_rate':[0.15,0.1,0.10,0.05], 'n_estimators':[100,150,200,250]}

tuning = GridSearchCV(estimator =GradientBoostingRegressor(),
                      param_grid = LR, scoring='r2')
tuning.fit(X_train,y_train)
```

```
from sklearn.model_selection import GridSearchCV
LR = {'learning_rate':[0.15,0.1,0.10,0.05], 'n_estimators':[100,150,200,250]}

tuning = GridSearchCV(estimator =GradientBoostingRegressor(),
                      param_grid = LR, scoring='r2')
tuning.fit(X_train,y_train)
tuning.best_params_, tuning.best_score_
```

```
({'learning_rate': 0.15, 'n_estimators': 200}, 0.8413790739479519)
```

Python